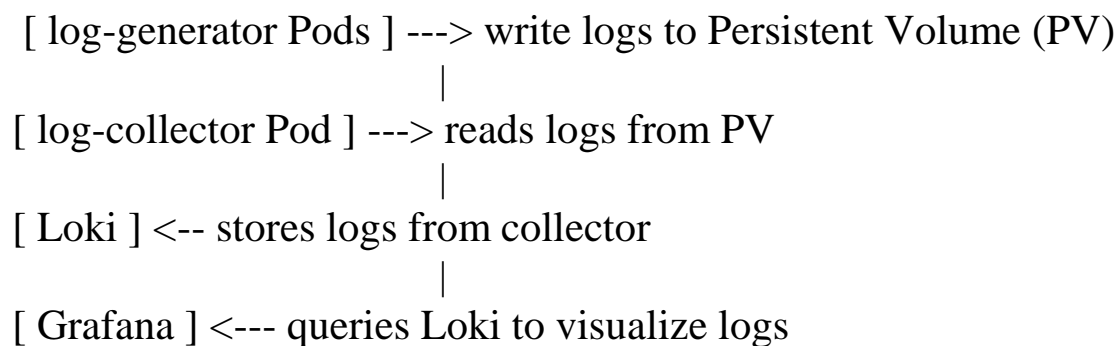


# Log Monitoring System

## Project Overview

- **Log Generator:** Simulates app logs
- **Log Collector:** Collects logs from mounted volumes (basic Python script or Filebeat)
- **Visualization:** Grafana or Kibana
- **Deployment:** Automated via shell scripts
- **Kubernetes:** Manages Pods, Services, Volumes

## Architecture Overview:



## Step 1: Create Directory Structure

```
mkdir -p log-monitoring-devops/{docker/log-generator,docker/log-collector,k8s,scripts}
cd log-monitoring-devops
```

## Step 2: Create Log Collector (Python-based)

```
docker/log-collector/collector.py
```

```
import time
```

```
import os
```

```
log_path = "/logs/access.log"
```

```
while True:
```

```
if os.path.exists(log_path):
    with open(log_path, "r") as f:
        lines = f.readlines()
        print("".join(lines[-10:]))
else:
    print("Waiting for logs...")
    time.sleep(5)
```

### [docker/log-collector/Dockerfile](#)

```
FROM python:3.9-slim
WORKDIR /app
COPY collector.py .
VOLUME ["/logs"]
CMD ["python", "collector.py"]
```

## Step 3: Build Docker Image

### [scripts/build.sh](#)

```
#!/bin/bash
eval $(minikube docker-env)
```

```
docker build -t log-collector:latest ./docker/log-collector
```

```
echo "Built log-collector image"
```

Run it:

```
chmod +x scripts/build.sh
./scripts/build.sh
```

## Step 4: Create Persistent Volume & PVC

### [k8s/pv.yaml](#)

```
apiVersion: v1
```

kind: PersistentVolume

metadata:

name: log-pv

spec:

capacity:

storage: 1Gi

accessModes:

- ReadWriteMany

hostPath:

path: /tmp/logdata

---

apiVersion: v1

kind: PersistentVolumeClaim

metadata:

name: log-pvc

spec:

accessModes:

- ReadWriteMany

resources:

requests:

storage: 1Gi

## Step 5: NGINX Log Generator Deployment

[k8s/log-generator.yaml](#)

apiVersion: apps/v1

kind: Deployment

metadata:

name: log-generator

spec:

replicas: 2

selector:

matchLabels:

app: log-generator

template:

metadata:

labels:

app: log-generator

spec:

containers:

- name: nginx  
image: nginx:alpine  
volumeMounts:
  - name: log-volume  
mountPath: /var/log/nginx

volumes:

- name: log-volume  
persistentVolumeClaim:  
claimName: log-pvc

## Step 6: Log Collector Pod

[k8s/log-collector.yaml](#)

apiVersion: v1

kind: Pod

metadata:

name: log-collector

spec:

containers:

- name: collector  
image: log-collector:latest  
volumeMounts:
  - name: log-volume  
mountPath: /logs

volumes:

- name: log-volume  
persistentVolumeClaim:  
claimName: log-pvc

## Step 7: Apply All YAMLs

[scripts/deploy.sh](#)

#!/bin/bash

kubectl apply -f k8s/pv.yaml

kubectl apply -f k8s/log-generator.yaml

```
kubectl apply -f k8s/log-collector.yaml
```

```
echo "Deployed PV, NGINX, and Collector"
```

```
kubectl get pods
```

Run it:

```
chmod +x scripts/deploy.sh
```

```
./scripts/deploy.sh
```

Check logs:

```
kubectl logs log-collector
```

## Step 8: Deploy Grafana + Loki (via Helm)

### Install Helm :

Install Helm for Windows

- Go to: <https://github.com/helm/helm/releases/latest>
- Download helm-v3.x.x-windows-amd64.zip
- Extract helm.exe and put it in a folder like C:\bin
- Add C:\bin to your system PATH

Then verify:

```
helm version
```

### Next Steps: Install Loki + Grafana in Minikube

1. Add the Grafana Helm repo:

```
helm repo add grafana https://grafana.github.io/helm-charts
```

```
helm repo update
```

2. Install Loki Stack with Grafana and Promtail enabled:

```
helm install loki grafana/loki-stack \
--set grafana.enabled=true \
--set promtail.enabled=true
```

This will install:

- Grafana (visualization)
- Loki (log aggregation)
- Promtail (log forwarder from your pods)

3. Check the services:

```
kubectl get svc
```

Look for a service like loki-grafana or similar.

4. Expose Grafana to view in browser:

```
minikube service loki-grafana --url
```

Open the URL in your browser — default credentials:

- Username: admin
- Password: prom-operator (or check using: `kubectl get secret --namespace default loki-grafana -o jsonpath="{.data.admin-password}" | base64 --decode`)

## Confirm log generation

If you already deployed your Flask app or log generator pods, Promtail should be picking up logs and sending them to Loki.

In Grafana:

1. Go to “Explore”
2. Select the “Loki” data source
3. Try this query:

```
{app="log-generator"}
```

You should start seeing logs appear!

## Important Notes

- Promtail must be deployed as DaemonSet and must access container logs.
  - Loki and Promtail pods must be healthy (kubectl get pods).
  - If you stop/start Minikube, Helm installations persist, but you must re-run minikube service to re-expose services.
- 

## Validation

- Logs appear in Grafana
  - Pods are running and generating logs
  - Grafana dashboard is accessible via browser
- 

## Optional Enhancements

- Add log-collector pod
- Set up Persistent Volumes
- Use Ingress for stable access to Grafana
- Add alerting in Grafana based on log patterns